

skin cancer detection

October 26, 2025

```
[ ]: import os
import numpy as np

import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
    Dropout, BatchNormalization, GlobalAveragePooling2D
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau

from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

[2]: BASE_DIR = "C:\\skin cancer detection" # Change based on your dataset location
TRAIN_DIR = os.path.join(BASE_DIR, "train")
TEST_DIR = os.path.join(BASE_DIR, "test")

# Image size and batch size
IMG_SIZE = (224, 224) # Optimal size for EfficientNetB0
BATCH_SIZE = 32

# Data Augmentation for Training
train_datagen = ImageDataGenerator(
    rescale=1./255, # Normalize pixel values to [0, 1]
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True
)

# Rescaling only for Testing
test_datagen = ImageDataGenerator(rescale=1./255)

# Load Training Data
```

```

train_generator = train_datagen.flow_from_directory(
    directory=TRAIN_DIR,
    target_size=IMG_SIZE,
    batch_size=BATCH_SIZE,
    class_mode="binary" # Binary classification: Benign vs Malignant
)

# Load Testing Data
test_generator = test_datagen.flow_from_directory(
    directory=TEST_DIR,
    target_size=IMG_SIZE,
    batch_size=BATCH_SIZE,
    class_mode="binary"
)

# Print class labels
print("Class Mapping:", train_generator.class_indices)

```

Found 11879 images belonging to 2 classes.

Found 2000 images belonging to 2 classes.


Class Mapping: {'Benign': 0, 'Malignant': 1}

```

[3]: # Display some sample images from dataset
def plot_images(generator):
    images, labels = next(generator) # Get a batch
    plt.figure(figsize=(10, 10))
    for i in range(9):
        plt.subplot(3, 3, i+1)
        plt.imshow(images[i])
        plt.title("Cancer" if labels[i] == 1 else "Non-Cancer")
        plt.axis("off")
    plt.show()

# Plot Training Images
plot_images(train_generator)

```



```
[4]: from tensorflow.keras import layers, models

# Build the CNN model
model = models.Sequential([
    # 1st Convolution Block
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(224, 224, 3)),
    layers.MaxPooling2D((2, 2)),

    # 2nd Convolution Block
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
```

```

# 3rd Convolution Block
layers.Conv2D(128, (3, 3), activation='relu'),
layers.MaxPooling2D((2, 2)),

# Flattening Layer
layers.Flatten(),

# Fully Connected Layers
layers.Dense(512, activation='relu'),
layers.Dropout(0.5), # Dropout to reduce overfitting
layers.Dense(1, activation='sigmoid') # Output layer for binary
↳classification
])

# Compile the model
model.compile(optimizer='adam', # Optimizer for faster convergence
              loss='binary_crossentropy', # Binary classification loss
              metrics=['accuracy']) # Accuracy metric

# Model Summary
model.summary()

```

C:\Users\adkel\anaconda3\Lib\site-packages\keras\src\layers\convolutional\base_conv.py:113: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 222, 222, 32)	896
max_pooling2d (MaxPooling2D)	(None, 111, 111, 32)	0
conv2d_1 (Conv2D)	(None, 109, 109, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 54, 54, 64)	0
conv2d_2 (Conv2D)	(None, 52, 52, 128)	73,856
max_pooling2d_2 (MaxPooling2D)	(None, 26, 26, 128)	0
flatten (Flatten)	(None, 86528)	0

dense (Dense)	(None, 512)	44,302,848
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 1)	513

Total params: 44,396,609 (169.36 MB)

Trainable params: 44,396,609 (169.36 MB)

Non-trainable params: 0 (0.00 B)

```
[5]: # EarlyStopping to stop training when the model stops improving
early_stopping = EarlyStopping(
    monitor="val_loss",
    patience=5, # Stop training after 5 epochs with no improvement
    restore_best_weights=True,
    verbose=1
)

# ReduceLROnPlateau to reduce the learning rate if the model stops improving
lr_scheduler = ReduceLROnPlateau(
    monitor="val_loss",
    factor=0.3,
    patience=3, # Reduce LR if no improvement after 3 epochs
    min_lr=1e-7,
    verbose=1
)

history = model.fit(
    train_generator,
    steps_per_epoch=train_generator.samples // BATCH_SIZE,
    epochs=30, # Set number of epochs to train the model
    validation_data=test_generator,
    validation_steps=test_generator.samples // BATCH_SIZE,
    callbacks=[early_stopping, lr_scheduler]
)
```

C:\Users\adkel\anaconda3\Lib\site-packages\keras\src\trainers\data_adapters\py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.

```

self._warn_if_super_not_called()

Epoch 1/30
371/371          4194s 11s/step -
accuracy: 0.7593 - loss: 0.5132 - val_accuracy: 0.7984 - val_loss: 0.4473 -
learning_rate: 0.0010
Epoch 2/30
 1/371          24:43 4s/step - accuracy:
0.8438 - loss: 0.3878

C:\Users\adkel\anaconda3\Lib\site-
packages\keras\src\trainers\epoch_iterator.py:116: UserWarning: Your input ran
out of data; interrupting training. Make sure that your dataset or generator can
generate at least `steps_per_epoch * epochs` batches. You may need to use the
`.repeat()` function when building your dataset.
  self._interrupted_warning()

371/371          55s 138ms/step -
accuracy: 0.8438 - loss: 0.3878 - val_accuracy: 0.7913 - val_loss: 0.4508 -
learning_rate: 0.0010
Epoch 3/30
371/371          1829s 5s/step -
accuracy: 0.8271 - loss: 0.3965 - val_accuracy: 0.8659 - val_loss: 0.3753 -
learning_rate: 0.0010
Epoch 4/30
371/371          52s 130ms/step -
accuracy: 0.9062 - loss: 0.3144 - val_accuracy: 0.8604 - val_loss: 0.3865 -
learning_rate: 0.0010
Epoch 5/30
371/371          1863s 5s/step -
accuracy: 0.8356 - loss: 0.3791 - val_accuracy: 0.8488 - val_loss: 0.3276 -
learning_rate: 0.0010
Epoch 6/30
371/371          186s 494ms/step -
accuracy: 0.7812 - loss: 0.4253 - val_accuracy: 0.8427 - val_loss: 0.3378 -
learning_rate: 0.0010
Epoch 7/30
371/371          7897s 21s/step -
accuracy: 0.8351 - loss: 0.3710 - val_accuracy: 0.8387 - val_loss: 0.3579 -
learning_rate: 0.0010
Epoch 8/30
 1/371          12:33 2s/step - accuracy:
0.7188 - loss: 0.5647
Epoch 8: ReduceLROnPlateau reducing learning rate to 0.0003000000142492354.
371/371          41s 105ms/step -
accuracy: 0.7188 - loss: 0.5647 - val_accuracy: 0.8463 - val_loss: 0.3497 -
learning_rate: 0.0010
Epoch 9/30
371/371          933s 3s/step -

```

```

accuracy: 0.8478 - loss: 0.3451 - val_accuracy: 0.8876 - val_loss: 0.2933 -
learning_rate: 3.0000e-04
Epoch 10/30
371/371          31s 79ms/step -
accuracy: 0.8125 - loss: 0.4997 - val_accuracy: 0.8846 - val_loss: 0.2971 -
learning_rate: 3.0000e-04
Epoch 11/30
371/371          938s 3s/step -
accuracy: 0.8535 - loss: 0.3396 - val_accuracy: 0.8503 - val_loss: 0.3385 -
learning_rate: 3.0000e-04
Epoch 12/30
  1/371          12:59 2s/step - accuracy:
0.9062 - loss: 0.2965
Epoch 12: ReduceLROnPlateau reducing learning rate to 9.000000427477062e-05.
371/371          32s 80ms/step -
accuracy: 0.9062 - loss: 0.2965 - val_accuracy: 0.8478 - val_loss: 0.3383 -
learning_rate: 3.0000e-04
Epoch 13/30
371/371          922s 2s/step -
accuracy: 0.8578 - loss: 0.3291 - val_accuracy: 0.8735 - val_loss: 0.3172 -
learning_rate: 9.0000e-05
Epoch 14/30
371/371          30s 76ms/step -
accuracy: 0.9375 - loss: 0.2666 - val_accuracy: 0.8760 - val_loss: 0.3156 -
learning_rate: 9.0000e-05
Epoch 14: early stopping
Restoring model weights from the end of the best epoch: 9.

```

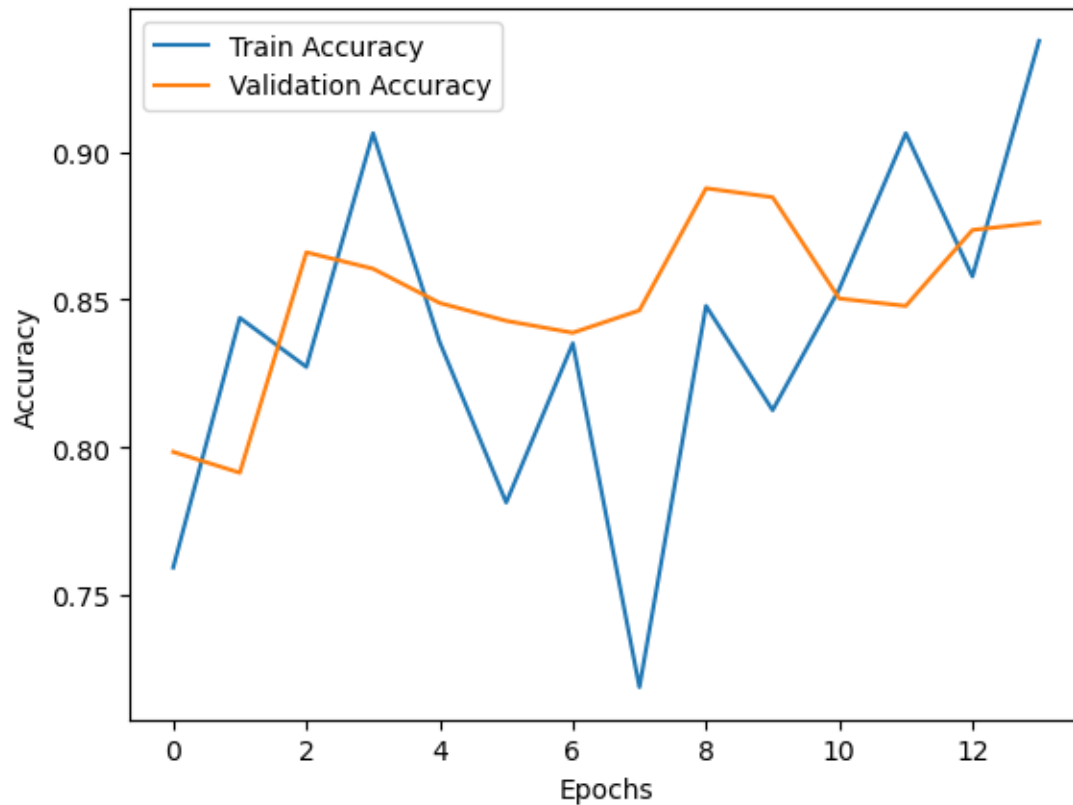
```

[7]: model.save("skin_cancer_cnn.h5")

# Plot Training History
plt.plot(history.history["accuracy"], label="Train Accuracy")
plt.plot(history.history["val_accuracy"], label="Validation Accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()

```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.



```
[8]: # Predict on the test set
test_pred = model.predict(test_generator, steps=test_generator.samples // ␣
    ↪ BATCH_SIZE, verbose=1)

# Convert predictions to binary labels (0 or 1)
test_pred_labels = (test_pred > 0.5).astype("int32")

# Get the true labels
test_true_labels = test_generator.classes[:len(test_pred_labels)] # Match␣
    ↪ length to predictions

# Classification report
print("Classification Report:")
print(classification_report(test_true_labels, test_pred_labels))

# Confusion Matrix
cm = confusion_matrix(test_true_labels, test_pred_labels)

# Plot Confusion Matrix
plt.figure(figsize=(6, 6))
```



```

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=train_generator.
    class_indices.keys(), yticklabels=train_generator.class_indices.keys())
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()

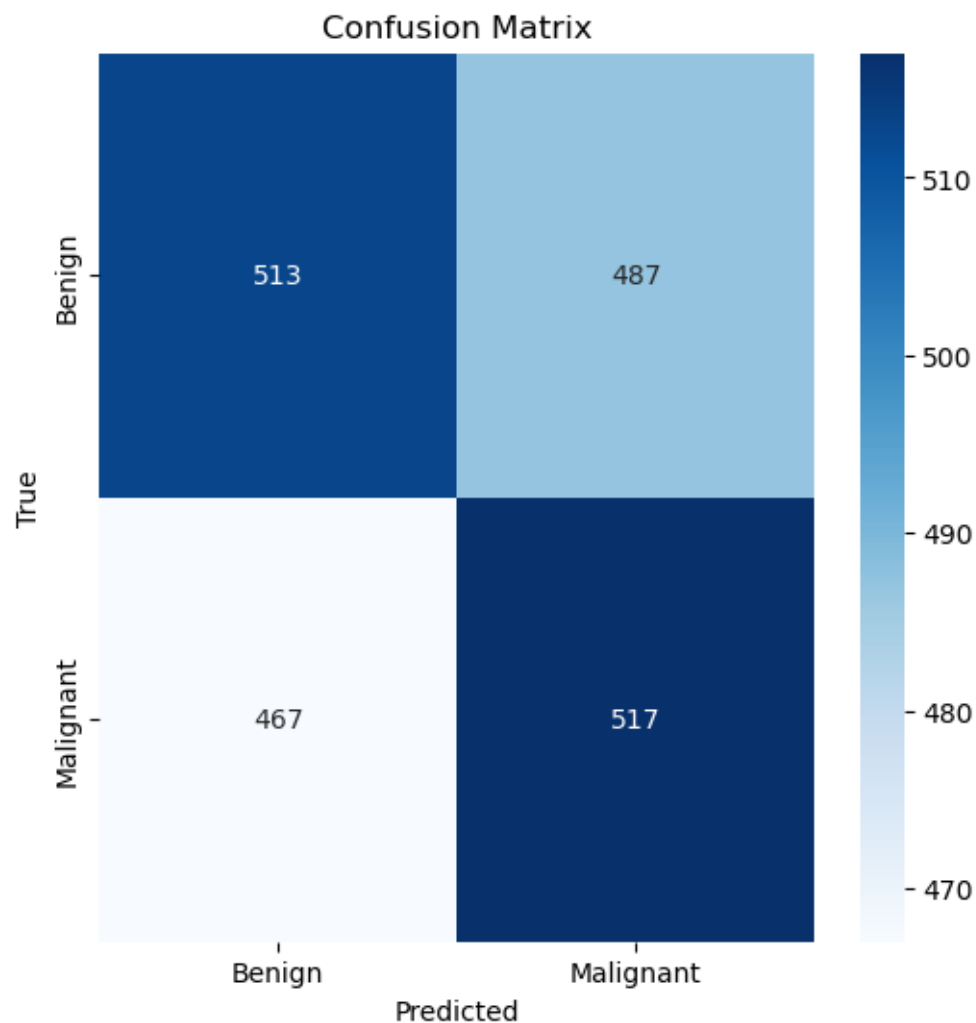
```

62/62

50s 816ms/step

Classification Report:

	precision	recall	f1-score	support
0	0.52	0.51	0.52	1000
1	0.51	0.53	0.52	984
accuracy			0.52	1984
macro avg	0.52	0.52	0.52	1984
weighted avg	0.52	0.52	0.52	1984



```
[13]: import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.preprocessing import image
from tensorflow.keras.models import load_model

# Load the entire model
model = load_model('skin_cancer_cnn.h5', compile=False)

def predict_skin_cancer(image_path, model):
    img = image.load_img(image_path, target_size=(224, 224)) # Load Image
    img_array = image.img_to_array(img) / 255.0 # Normalize
    img_array = np.expand_dims(img_array, axis=0) # Add batch dimension

    # Make Prediction
    prediction = model.predict(img_array)
    class_label = "Malignant" if prediction > 0.5 else "Benign"

    # Show Image with Prediction
    plt.imshow(img)
    plt.title(f"Predicted: {class_label}")
    plt.axis("off")
    plt.show()

[17]: predict_skin_cancer(r"C:\skin cancer detection\test\Benign\6331.jpg", model)
```

1/1 0s 115ms/step

Predicted: Benign



```
[19]: predict_skin_cancer(r"C:\skin cancer detection\train\Malignant\386.jpg", model)
```

1/1

0s 123ms/step

Predicted: Malignant



[]: